

Min-Han Li

+1-781-475-7402 | minhanli.wesley@gmail.com | [linkedin.com/in/minhanliwesley](https://www.linkedin.com/in/minhanliwesley) | San Jose, CA

EDUCATION

Northeastern University

Boston, MA

Master of Science in Computer Science, Minor in Digital Media

09/2024 – 04/2026

Coursework: Building Scalable Distributed Systems, Cloud Computing, Object-Oriented Design, Software Development

National Taiwan University

Taipei, Taiwan

Bachelor of Science, Double Major in Computer Science and Chemical Engineering

09/2016 – 05/2022

Coursework: Advanced Web Development, Data Processing and Analysis, Data Structures and Algorithms

EXPERIENCE

Software Engineer Intern

05/2025 – 11/2025

ByteDance (TikTok) — Go, Python, SQL, Redis, Docker, RAG, Multi-Agent Systems

San Jose, CA

- Shipped AI-powered on-call diagnostic bot in Go serving 200+ internal users, reducing average resolution time 50% (7→3.5 days) by auto-aggregating ads data and relevant past incidents at ticket creation
- Enabled knowledge-base-driven self-service powered by RAG and LLM-based routing that automates 50% of 30+ weekly ad-platform issues, eliminating dependency on on-call engineers for common incidents
- Developed 6+ integrated troubleshooting tools in Go — permission verification, access checks, log-based user lookup, and one-click bug filing — behind a unified card interface accessible to non-technical staff
- Designed multi-agent diagnostic workflows from internal runbooks for ad delivery debugging, reducing manual escalations previously involving 10+ people per incident
- Built knowledge pipeline converting resolved incidents into searchable articles via LLM chat summarization and RAG, with Redis caching on high-frequency lookups to minimize query latency

Software Engineer Intern

11/2024 – 01/2025

ReMo — Python, Go, FastAPI, YOLOv11, Docker, React, Zustand, AWS ECS

Portland, ME

- Developed backend services using Golang and FastAPI, enabling sub-3s inference latency for multi-stage OCR and metadata extraction
- Engineered a 90% accurate book spine recognition pipeline with YOLOv11, Internvl, and custom LLM-based metadata extraction
- Optimized frontend bundle size and rendering performance by implementing lazy-loading and code-splitting in React, reducing initial page load time by 25%
- Containerized and deployed services via Docker, AWS ECS, and API Gateway for high availability

Software Engineer Intern

07/2023 – 08/2023

Yanray Inc. — C#, .NET Core, Azure, Redis, Cosmos DB, Robot Framework

Taipei, Taiwan

- Delivered 95% structured field extraction accuracy by building a business card recognition pipeline with Azure AI Vision and a custom LLM
- Boosted metadata retrieval speed by 31% (175ms → 120ms) by caching reference lookups in Redis and optimizing Cosmos DB queries
- Applied SDLC best practices to develop, test, and deploy the backend on Azure App Service, integrating telemetry with Application Insights

Research & Development Intern

07/2022 – 08/2022

Acer — Python, OpenVINO, YOLOv7, FaceNet, Model Quantization, PyTorch, ONNX

Taipei, Taiwan

- Built a real-time meeting assistant that performed face and QR recognition under 300ms latency by combining model training, OpenVINO quantization, and async inference
- Trained and fine-tuned YOLOv7 and FaceNet models for QR detection and face recognition, optimizing for both accuracy and deployment efficiency
- Reduced model latency by 50% through INT8 quantization and model conversion for CPU deployment

TECHNICAL SKILLS

Languages: Java, Python, C++, JavaScript, C#, TypeScript, Ruby, Golang, Swift, Scala

Frameworks & Tools: Spring Boot, .NET, Dropwizard, Node.js, React, Maven, JUnit, Tomcat, Bazel

DevOps: AWS, GCP, Azure, Kubernetes, Message Queue, OpenFeign, Eureka, Docker, Microservices, RabbitMQ

Databases: MySQL, PostgreSQL, Redis, MongoDB, NoSQL, Iceberg, MyBatis, Hibernate, Spark, Hadoop

Machine Learning: PyTorch, XGBoost, scikit-learn, TensorFlow **OS:** macOS, Windows, Linux