

Min-Han Li

+1-781-475-7402 | minhanli.wesley@gmail.com | [linkedin.com/in/minhanliwesley](https://www.linkedin.com/in/minhanliwesley) | San Jose, CA

EDUCATION

Northeastern University

Boston, MA

Master of Science in Computer Science, Minor in Digital Media

09/2024 – 04/2026

Coursework: Building Scalable Distributed Systems, Cloud Computing, Object-Oriented Design, Software Development

National Taiwan University

Taipei, Taiwan

Bachelor of Science, Double Major in Computer Science and Chemical Engineering

09/2016 – 05/2022

Coursework: Advanced Web Development, Data Processing and Analysis, Data Structures and Algorithms

EXPERIENCE

Software Engineer Intern

05/2025 – 11/2025

ByteDance (TikTok) — Go, Python, SQL, Redis, Docker, RAG, Multi-Agent Systems

San Jose, CA

- Shipped production **backend service in Go** serving **200+ internal users**, reducing average incident resolution time **50%** (7→3.5 days) by auto-aggregating ads data and relevant past incidents at ticket creation
- Owned **end-to-end delivery** of 6+ integrated tools — permission verification, access checks, log-based user lookup, one-click bug filing — behind a unified card interface accessible to non-technical staff
- Built knowledge pipeline converting resolved incidents into searchable articles via **LLM** summarization and **RAG**, with **Redis** caching on high-frequency lookups to minimize query latency
- Designed multi-agent diagnostic workflows from internal runbooks, reducing manual escalations that previously required **10+ people** per incident
- Enabled knowledge-base-driven self-service powered by **RAG** and **LLM-based routing**, automating **50% of 30+ weekly ad-platform issues** and eliminating dependency on on-call engineers for common incidents

Software Engineer Intern

11/2024 – 01/2025

ReMo — Python, Go, FastAPI, YOLOv11, Docker, React, Zustand, AWS ECS

Portland, ME

- Optimized **React** frontend with lazy-loading and code-splitting, reducing initial page load time by **25%**
- Developed **Python + Go + FastAPI** backend services achieving **sub-3s inference latency** for multi-stage OCR and LLM-based metadata extraction
- Engineered a **90% accurate** book-spine recognition pipeline with **YOLOv11**, Internvl, and custom **LLM** metadata extraction
- Containerized and deployed services via **Docker**, **AWS ECS**, and API Gateway for high availability

Software Engineer Intern

07/2023 – 08/2023

Yanray Inc. — C#, .NET Core, Azure, Redis, Cosmos DB, Robot Framework

Taipei, Taiwan

- Delivered **95% structured field extraction accuracy** by building a business-card recognition pipeline with **Azure AI Vision** and a custom **LLM**
- Boosted metadata retrieval speed **31%** (175ms→120ms) by caching reference lookups in **Redis** and optimizing Cosmos DB queries
- Applied SDLC best practices to develop, test, and deploy the backend on **Azure App Service**, integrating telemetry with Application Insights

Research & Development Intern

07/2022 – 08/2022

Acer — Python, OpenVINO, YOLOv7, FaceNet, Model Quantization, PyTorch, ONNX

Taipei, Taiwan

- Built a real-time meeting assistant performing face and QR recognition under **300ms latency** by combining model training, **OpenVINO** quantization, and async inference
- Trained and fine-tuned **YOLOv7** and **FaceNet** for QR detection and face recognition, optimizing accuracy and deployment efficiency
- Reduced model latency **50%** through **INT8 quantization** and CPU-targeted model conversion

TECHNICAL SKILLS

Languages: TypeScript, JavaScript, Go, Python, Java, C++

Frontend: React, Zustand, Next.js, TypeScript, responsive design

Backend: Node.js, FastAPI, Spring Boot, gRPC, REST APIs, GraphQL, microservices

Distributed Systems: Redis, PostgreSQL, MySQL, MongoDB, queuing, caching, retries, reliability patterns

AI / ML: RAG, LLM integration, multi-agent systems, PyTorch, YOLO, OpenVINO **DevOps / Cloud:** AWS (ECS, API Gateway), GCP, Azure, Docker, Kubernetes, CI/CD