

Min-Han Li

+1-781-475-7402 | minhan.li.wesley@gmail.com | [LinkedIn](#) | [GitHub](#) | San Jose, CA

EDUCATION

Northeastern University

Boston, MA

Master of Science in Computer Science, Minor in Digital Media

09/2024 – 04/2026

- Coursework: Building Scalable Distributed Systems, Cloud Computing, Object-Oriented Design

National Taiwan University

Taipei, Taiwan

B.S., Double Major in Computer Science and Chemical Engineering

09/2016 – 05/2022

EXPERIENCE

Software Engineer Intern

05/2025 – 11/2025

ByteDance (TikTok) | Python, Go, RAG, Multi-Agent, LLM Tools, Redis

San Jose, CA

- Shipped AI-powered on-call diagnostic bot in Go serving **200+** internal users, reducing average resolution time **50%** (7→3.5 days) by auto-aggregating data and retrieving past incidents at ticket creation
- Enabled knowledge-base-driven self-service powered by RAG and LLM-based routing that automates **50%** of 30+ weekly issues, eliminating dependency on on-call engineers for common incidents
- Developed **5+** Go services that the diagnostic LLM agent invokes as tools for permission verification, access checks, log lookup, and bug filing, behind a unified card interface
- Designed multi-agent diagnostic workflows from internal runbooks, reducing manual escalations previously involving **10+** specialists per incident
- Built knowledge pipeline converting resolved incidents into searchable articles via LLM chat summarization and RAG, with Redis caching to minimize query latency

Software Engineer Intern

11/2024 – 01/2025

ReMo | Python, Go, FastAPI, YOLOv11, React, AWS ECS | [GitHub](#)

Portland, ME

- Delivered ReMo's MVP for fundraising: a **90%** accurate book spine recognition pipeline combining YOLOv11 detection, Internvl recognition, and custom LLM-based metadata extraction
- Achieved **sub-3s** end-to-end inference latency for the multi-stage OCR and metadata extraction system by building Python and Go backend services with FastAPI
- Optimized React frontend bundle size and rendering performance with lazy-loading and code-splitting, reducing initial page load time by **25%**
- Containerized and deployed services via Docker, AWS ECS, and API Gateway for high availability

Research & Development Intern

07/2022 – 08/2022

Acer | Python, PyTorch, OpenVINO, ONNX, YOLOv7, FaceNet

Taipei, Taiwan

- Built a real-time meeting assistant POC in Python performing face and QR recognition under **300ms** latency for up to **6 concurrent users**, combining PyTorch model training, OpenVINO quantization, and async inference
- Reduced model latency by **50%** through INT8 quantization and ONNX conversion for CPU deployment
- Submitted as a **patent** for hardware-level integration in Acer laptops, with YOLOv7 and FaceNet models optimized for on-device CPU inference

Machine Learning Engineer Intern

01/2021 – 06/2021

Formosa Plastics | Python, React, Flask, PyTorch, AutoML, Cantera

Taipei, Taiwan

- Shipped a full-stack ML product end-to-end for Formosa engineers designing a new production facility, owning data pipeline, model training, Flask backend, and React frontend
- Reduced model prediction error **75%** and accelerated the deployed pipeline **3×** over the prior toolchain by training an AutoML-tuned residual neural network on **80K** production measurements

PUBLICATIONS

B.-C. Ke, **M.-H. Li**, et al. *TurnAhead: Designing 3-DoF Rotational Haptic Cues to Improve First-person Viewing (FPV) Experiences*. **CHI 2023, Honorable Mention Award**. Prototyped a VR haptic feedback system in Unity / Meta Quest 2 / Arduino with 3D-printed hardware; validated through 4 user studies with 30+ participants. doi:10.1145/3544548.3581443

TECHNICAL SKILLS

AI & ML: PyTorch, RAG, LLM tooling, agentic systems, multi-agent workflows, scikit-learn, XGBoost, TensorFlow, ONNX, OpenVINO, quantization

Python & Backend: Python, FastAPI, Flask, Go, Docker, Kubernetes, Git, AWS, GCP, Azure, Redis, PostgreSQL, MySQL, MongoDB, Spark

Frontend & Other: TypeScript, JavaScript, React, Java, C++, Scala